

## Overview and Challenges

Indian Railways train scheduling is a **large-scale, dynamic combinatorial problem**. Dispatchers must manually sequence heterogeneous trains (express, freight, local, etc.) on limited track sections and platforms, respecting safety rules and train priorities. As noted in practice, “train scheduling is a very large, complex, combinatorial problem” <sup>1</sup>, traditionally managed by experienced controllers. However, growing congestion and expectations for punctuality demand intelligent decision support. In particular, express trains (with fewer stops) are given explicit network priority <sup>2</sup>, and safety rules (block signaling, separation times, etc.) must never be violated. An automated system must therefore encode all operational constraints (track occupancy, safety blocks, platform limits, precedence rules, etc.) and **optimize throughput and delay** under these rules. Human dispatchers often use simple heuristics (e.g. delaying lower-priority trains) but may miss globally better schedules. Research shows that **automated decision-support systems** with optimization have proven effective in practice <sup>1</sup>. The goal is an AI-assisted controller: it suggests conflict-free, high-throughput train sequences while allowing overrides, and rapidly re-optimizes when disruptions occur.

## Data Sources and API Integration

Building such a system requires **rich, real-time data** on train movements, schedules, and infrastructure. Key sources include:

- **Official timetable and train data.** The Indian government’s Open Data portal provides static datasets. For example, the Ministry of Railways published a complete *Train Time Table* (train-wise schedules and routes) for all reserved trains (as of a given date) <sup>3</sup>. This can bootstrap the baseline schedule. (In practice, data.gov.in resources such as “Indian Railways Time Table” supply arrival/departure times for every station on each train’s route.)
- **Public APIs.** Several third-party APIs offer Indian Rail info:
  - *IndianRailAPI.com* provides endpoints to look up station codes, train routes, schedules, live running status, and PNR status <sup>4</sup>. For example, one can fetch “trains between stations” or the “current running status” via HTTP APIs <sup>4</sup>.
  - *eRail API* (api.erail.in) is a public service with real-time train data: schedules, fares, coach info, seat availability, etc. It covers ~7,500 trains and 22,000 stations across India <sup>5</sup>. Using an API key, one can query live train status, get the full timetable of a train, find trains between stations, etc. The eRail API’s documentation shows straightforward REST calls to retrieve schedules or live status in JSON <sup>5</sup>.
  - *RapidAPI/IRCTC* and other endpoints exist for live train status and search (these are unofficial and often require API keys).
- **Station and network data.** We need geography of stations and the network topology. Government datasets often **lack station coordinates**, so these must be supplemented. For example, Suratekar (2018) found he had to scrape station latitudes/longitudes from external sources because the official timetable data lacked geolocation <sup>6</sup>. In practice, one can use **OpenStreetMap/OpenRailwayMap** data for track layouts, station locations, and signal information. OpenRailwayMap (an OSM-based

project) provides a “world-wide, open, up-to-date and detailed map of the railway network” <sup>7</sup> .

Using its API or raw OSM data, one can construct the rail graph (track segments, junctions, signaling constraints).

- **Real-time updates.** For live operations, streaming updates of train positions and delays are needed. This can come from IRCTC live status APIs, third-party trackers (e.g. WhereIsMyTrain app APIs or crowdsourced feeds), or eventually from the Railway’s Train Management System (TMS) if accessible via secure feeds. These data feed into the system to trigger re-optimization on disturbances.

**Possible API integrations:** Real-time APIs (e.g. Live Train Status) should be used as a first principle. For example, a Rail API call can return the current delay or location of a given train. If official APIs are closed, third-party APIs like eRail or apiclub’s “Live Train Status” (often on RapidAPI) can be used (some have free tiers). Station information can be fetched from OSM or Google Maps Geocoding for missing coords. In summary, our data pipeline should **combine official open-data, third-party rail APIs, and open-source GIS data** to cover all needed inputs <sup>6</sup> <sup>5</sup> .

## Data Engineering Pipeline

The data engineering tasks include **ingesting, storing, and transforming** the above data for optimization:

- **ETL Ingestion:** Schedule regular extraction jobs (e.g. daily) to pull static timetables from data.gov.in or other sources. Use periodic API calls (cron or Airflow) to fetch dynamic data (live statuses, delays) during operations. Tools like Python’s `requests` / `aiohttp` libraries or Airflow DAGs can orchestrate this.
- **Data Cleaning and Integration:** Resolve differences in naming (e.g. station code vs. name). Suratekar noted (2018) that merging timetable data with station geodata required custom processing <sup>6</sup> . We must unify reference frames: all times to a common time zone, station codes from different sources to a single canonical set, etc. Missing data (e.g. station lat/long) must be filled via geocoding or external datasets.
- **Data Storage:** Use a robust database (e.g. PostgreSQL or time-series DB like TimescaleDB) to store schedules, real-time updates, track topology, and past performance. For real-time feeds (e.g. frequent train updates), a message queue (Kafka/RabbitMQ) or streaming DB can buffer incoming data. The database schema would include tables for trains (ID, class, priority), station network (station nodes, track edges, capacities), timetable entries (arrival/departure times), and dynamic events (delay reports).
- **Feature Computation:** Pre-compute derived data such as segment traversal times, average speeds, or conflict matrices (which trains share tracks). For machine learning modules, also prepare historical features (e.g. delay histories).
- **Orchestration:** Tools like Apache Airflow or Prefect can schedule pipeline steps (data fetch, preprocess, store). Alternatively, a custom Python ETL framework can suffice. The pipeline must handle failures (retry on API timeouts) and log its actions for auditing.

Throughout this pipeline, **security and compliance** matter: any integration with official systems must use secure APIs (HTTPS, proper auth). For example, if the system eventually connects to a TMS, we would use token-based or certificate-based auth. All data transformations should comply with IR data usage policies.

# Backend Architecture

The **backend** can be designed as a set of microservices or a modular monolith, depending on scale:

- **API/Service Layer:** A core service (e.g. Python Flask/FastAPI or Node.js) exposes endpoints for internal modules and possibly external UIs. Services include:
  - A *Scheduling Service* that accepts current state (incoming trains, track status, delays) and runs the optimization model (see below).
  - A *Data Service* that serves up stored timetables, network topology, and historical KPIs to other modules.
  - A *Simulation/Scenario Service* for what-if analyses (below).
- **Database:** As noted, a relational or graph DB stores the data. For graph analyses (finding shortest paths on the network), a graph database (Neo4j) or spatial DB (PostGIS) could help, though not strictly required.
- **Queue/Stream:** If handling real-time updates, a message queue (Kafka) can decouple data ingestion from processing. For example, each new train position or delay can publish to a topic that the Scheduling Service subscribes to trigger re-optimization.
- **UI / Web Dashboard:** Controllers need a user interface. This could be a web app (React/Angular) or a desktop interface. It displays the current schedule plan (timeline/graph view), alerts, and allows manual overrides. REST API endpoints from the backend provide the recommended next moves. The UI must clearly show **recommendations and rationale**, e.g. "Hold train A at station X (to let higher-priority train B pass) because track Y is occupied <sup>8</sup>."
- **Authentication & Security:** All external interfaces should use HTTPS with proper authentication. If integrating with IR's systems, likely OAuth2 or JWT tokens. Audit logging must record every decision sent to UI and every override by a controller.

A possible stack: Python for core logic (data ingestion, OR-Tools), PostgreSQL for storage, Redis for caching recent queries, and a frontend in a modern JS framework for UI. Containerization (Docker) and orchestration (Kubernetes) can make it scalable. Infrastructure monitoring (Prometheus/Grafana) will track system health.

## Optimization with OR-Tools

The **heart** is an optimization model implemented in Python using Google OR-Tools. We must encode:

- **Variables:** For each train and each track segment or station event, define start and end times. For example, variables  $t_{i,s}$  = time train  $i$  enters segment or station  $s$ .
- **Constraints:**
  - **Track occupancy:** No two trains can occupy the same track segment at overlapping times. For each segment shared by trains  $i$  and  $j$ , add disjunctive constraints (e.g.  $t_{i,\text{end}} \leq t_{j,\text{start}}$  OR vice versa). In CP-SAT this can use boolean indicator variables or NoOverlap constraints.
  - **Precedence/switching:** When trains cross at a junction or station, enforce safety separation. For instance, signals require a minimum clearance time between trains.
  - **Train order/preference:** Hard constraints encode safety; soft preferences encode priorities. For example, we allow a lower-priority train to be delayed if needed. We can model *train priorities* (e.g. Rajdhani > Garib Rath > freight) by adding weighted penalties: delaying a high-priority train adds a

large cost, so the optimizer favors letting it run on time. As noted, “express trains...having fewer halts [are given] priority on [the] rail network” <sup>2</sup>, so our objective will penalize delays for high-priority services more.

- **Platform availability:** Limit how many trains can occupy a station (or platform) concurrently.
- **Service time:** If a train stops at a station, its dwell time must meet minimums.
- **Rolling stock constraints** (optional): If considering locomotive or rake reuse, ensure one train’s schedule respects its equipment availability.
- **Objective:** Typically a multi-objective or weighted sum. Common objectives are **minimize total delay or maximize throughput**. For example, minimize  $\sum_i w_i \cdot \text{delay}_i$ , where  $w_i$  is weight based on priority. Another goal is to minimize the *makespan* or completion time of all trains through the section. Research often uses makespan or total completion time (as in freight scheduling) <sup>9</sup>. We could also include penalties for deviations from timetable to favor on-time performance.
- **Solving:** Use OR-Tools CP-SAT (Constraint Programming with SAT). CP-SAT is a state-of-the-art solver that can handle large combinatorial models. In one case study, an **exact optimal solution** for a rail line schedule was obtained using the OR-Tools solver <sup>9</sup>. We will configure the solver with reasonable time limits (e.g. 1–5 seconds for near-real-time response) and objective gap tolerances. For very large networks, solve incrementally (e.g. optimize one section at a time, then coordinate) or use heuristics to generate a good initial solution (OR-Tools supports search hints and warm starts).

In Python, the workflow is: define the CP model (binary and integer variables, add constraints, set objective), then call `solver.Solve()`. OR-Tools supports linear constraints, precedence, and custom interval constraints, which fit scheduling problems. This covers “leveraging operations research and AI to model constraints, train priorities, and operational rules” as required.

## AI/ML Components

Beyond pure optimization, machine learning can enhance decision support:

- **Delay Prediction:** Use historical data to train models (e.g. time-series, regression, or LSTM neural nets) that predict likely delays or travel times under varying conditions (weather, congestion). Such predictions can feed into the optimizer to anticipate disruption effects.
- **Conflict Resolution Advisor:** In disturbances, dispatchers often decide *which* train to hold/release. We can frame this as a learning problem: given a conflict (two trains wanting the same track), predict which choice leads to better network performance. For instance, Suratekar et al. (2021) categorize tactics like retiming, rerouting, reordering <sup>8</sup>. A machine-learned policy (e.g. a classifier or reinforcement learning agent) could suggest which train to prioritize. Indeed, the Swedish MATRIX project notes that AI/ML approaches are “often inevitable” for fast, intelligent disturbance handling <sup>10</sup>. We might train such a model from historical dispatch decisions or simulation.
- **Scenario Analysis (What-if):** We can use ML to generate realistic disruption scenarios (e.g. probabilistic delays) and then run optimization to see the impact. Alternatively, a reinforcement-learning agent could learn scheduling policies over many simulated scenarios. Research in train rescheduling uses deep RL to quickly adapt timetables after disruptions <sup>10</sup> <sup>11</sup>. Even if not fully deployed, these can provide insights.

- **User Recommendations Explanation:** To make the tool explainable, we can apply ML (or rule-based) to generate human-readable justifications (e.g. “Optimized throughput by delaying freight train A by 5 min” with reasoning). This is an active research area.

## Compliance and Safety

The system must **adhere to all Indian Railways standards**. This means:

- **Signaling and Block Rules:** Ensure no two trains violate the same-block occupancy. The model inherently enforces this, but the parameters (block length, signalling overlap) must match IR practice.
- **Regulatory Constraints:** Follow IR’s Operating Rules: e.g. avoid reversing movements without clearance, respect speed limits on gradients, etc. These can be added as constraints or checks.
- **Safety Priorities:** In India, passenger train safety is paramount. The system should never suggest a move that compromises safety for efficiency. This is enforced by hard constraints above.
- **Auditing:** Every recommended move and override must be logged. This supports later regulatory review and continuous improvement.
- **Standards Alignment:** The optimization logic can be validated against small real-world scenarios and past incidents to ensure it replicates legal precedence rules (for example, always stopping before a catch siding as per signal rules). If in doubt, institutional “traffic knowledge” should be encoded as fixed rules, not left to AI guesswork.

By construction, the solution space only includes *feasible* schedules that meet IR’s rulebook. We would also consult domain experts (Section Controllers) to validate that suggested schedules are operationally sensible.

## End-to-End Prototype Workflow

An **example workflow** (applied per corridor or section) is:

1. **Select Corridor.** For initial prototyping, choose a busy but bounded corridor (e.g. the Delhi–Mumbai trunk line or a suburban zone). This defines the relevant network nodes (stations, junctions) and scheduled trains.
2. **Load Static Data.** Fetch the train timetable dataset from data.gov.in (e.g. “Trains available for reservation” dataset <sup>3</sup>) and import it into the database. Ingest track geometry from OpenStreetMap (using its API or a downloaded map) to build the track graph <sup>7</sup>.
3. **Populate DB.** Using the data-engineering pipeline, parse and store schedules (each train’s station list and times), station coordinates, train classes, and track segments. Calculate nominal travel times between stops.
4. **Real-Time Feed Setup.** Subscribe to a real-time train status API. As trains run, ingest delay reports (e.g. train 12345 is +10min at Station X). Publish these into the system.
5. **Trigger Optimization.** Whenever there is a new disturbance (delay or incoming unscheduled train), send the updated state to the Scheduling Service. It fetches all *active trains in the section*, their target schedules, and the current time-space conflicts.

6. **Build OR-Tools Model.** Programmatically create variables for each train's planned departure/arrival at upcoming segments, impose all constraints (as above), and set the objective (minimize weighted delays).
7. **Solve.** Run the OR-Tools solver. If time permits, allow it to find an optimal or near-optimal solution (it may use up to a preset time limit). The output is a conflict-free schedule (new times and track assignments).
8. **Post-Process.** Compare the new plan to current plan; identify which trains will be held or rerouted. For example, the plan might say "hold Freight Train F at Junction Y for 5 minutes, let Express Train E pass first".
9. **UI Notification.** Display recommendations on the controller's dashboard. Include a textual explanation: e.g. "Train E has higher priority <sup>2</sup> and is delayed, so routing it first reduces total delay." The controller can accept or override.
10. **Logging and KPIs.** Record the decision and track key metrics: how much delay was avoided, throughput in trains/hour, platform utilization, etc. Over time, update performance dashboards (e.g., Grafana charts of on-time performance).
11. **What-If Simulations.** In parallel, allow off-line scenario testing: e.g. "if an extra special train comes at 6pm, how to adjust the schedule?" This uses the same model with altered inputs. Simulation results can be exported for review.

Throughout, the system must be extensible to multiple corridors. After initial validation, we could scale it zone-by-zone: each Section Controller's area runs its own optimizer using local data, but a central service could coordinate overlaps.

**Technology Stack (example):** Python 3 backend using OR-Tools (CP-SAT), PostgreSQL/PostGIS for data, Redis for cache, Kafka for real-time streams, and a React-based web UI. Data ingestion and nightly batch jobs orchestrated via Apache Airflow. Monitoring via Prometheus/Grafana. All services containerized with Docker.

## References

- Official Indian Railways open timetable dataset <sup>3</sup> .
- Third-party APIs (e.g. IndianRailAPI.com, eRail API) <sup>4</sup> <sup>5</sup> .
- Academic studies on train scheduling and decision-support, e.g. freight train optimization using OR-Tools <sup>9</sup> , and ML-driven rescheduling (MATRIX project) <sup>12</sup> <sup>10</sup> .
- OpenStreetMap/OpenRailwayMap data for rail topology <sup>7</sup> .
- Indian Railways operational guidelines (e.g. train priority rules) <sup>2</sup> .

These sources inform the design of the data feeds, optimization logic, and system requirements for the end-to-end solution.

---

<sup>1</sup> <sup>9</sup> [scitepress.org](https://www.scitepress.org/PublishedPapers/2022/113795/113795.pdf)  
<https://www.scitepress.org/PublishedPapers/2022/113795/113795.pdf>

<sup>2</sup> [Indian Railways - Wikipedia](https://en.wikipedia.org/wiki/Indian_Railways)  
[https://en.wikipedia.org/wiki/Indian\\_Railways](https://en.wikipedia.org/wiki/Indian_Railways)

3 6 Indian Railways part I - travel to the Moon

<https://weirddata.github.io/2018/07/25/railway-stats.html>

4 API - INDIAN RAIL API

<https://indianrailapi.com/api-collection>

5 Indian Railways API — Free Public API | Public APIs Directory

<https://publicapis.io/indian-railways-api>

7 OpenRailwayMap - OpenStreetMap Wiki

<https://wiki.openstreetmap.org/wiki/OpenRailwayMap>

8 10 12 fudinfo.trafikverket.se

[https://fudinfo.trafikverket.se/fudinfoexternwebb/Publikationer/Publikationer\\_007701\\_007800/Publikation\\_007715/MATRIX\\_final\\_report.pdf](https://fudinfo.trafikverket.se/fudinfoexternwebb/Publikationer/Publikationer_007701_007800/Publikation_007715/MATRIX_final_report.pdf)

11 Reinforcement Learning for Scalable Train Timetable Rescheduling ...

<https://arxiv.org/html/2401.06952v1>